

CS 498 AMO (MCS-DS online) – Homework 1

Submitted by: Pushpit Saxena (netid: pushpit2)

Part 1 Accuracies

Setup	Cross-validation Accuracy
Unprocessed data	75.620915% *
0-value elements ignored	75.032680% *

* The accuracy mentioned here is from one of the 10 test-train split run. On average I see the accuracy for either of the case hovering in the range of (73.5% - 77%) for different runs.

Part 1 Code Snippets

Calculation of distribution parameters (Page 6 - Page 7)

```
1. def fit(self, X, Y):
2.     self.normDF = {}
3.     self.priors = {}
4.     categories = set(Y)
5.     if self.ignoreMissingVal:
6.         X[X == 0] = np.nan
7.     for c in categories:
8.         XForC = X[Y == c]
9.         self.normDF[c] = {
10.             'mean': np.nanmean(XForC, axis=0),
11.             'var': np.nanvar(XForC, axis=0)
12.         }
13.         self.priors[c] = 1.0 * len(Y[Y == c]) / len(Y)
```

Calculation of naive Bayes predictions (Page 7 - Page 8)

```
1. def predict(self, X):
2.     P = {}
3.     for c, g in self.normDF.items():
4.         mean, var = g['mean'], g['var']
5.         classConditionalProb = 0
6.         # Please see the detailed method implementation in the full code attached a
7.         classConditionalProb = np.sum(self.__calculateLogNormPdf(X, mean, np.sqrt(v
8.         P[c] = classConditionalProb + np.log(self.priors[c])
9.         bestCategory, bestProb = None, float("-inf")
10.     for category, probability in P.items():
11.         if bestCategory is None or probability > bestProb:
12.             bestProb = probability
13.             bestCategory = category
14.     return bestCategory
```











Test-train split code (Page 6)

```
1. def testTrainSplit(self, data, ratio):
2.     localCopy = list(data)
3.     testSize = int(len(data) * ratio)
4.     testData = []
5.     while len(testData) < testSize:
6.         testData.append(localCopy.pop(random.randrange(len(localCopy))))
7.     testNPArr = np.array(testData)
8.     trainNPArr = np.array(localCopy)
9.     return trainNPArr[:, :8], trainNPArr[:, 8].astype(int), testNPArr[:, :8], testNPArr[:, 8].astype(int)
```

Part 2 MNIST Accuracies

x	Method	Training Set Accuracy	Test Set Accuracy
1	Gaussian + untouched	79.041667	79.860000
2	Gaussian + stretched	81.656667	82.450000
3	Bernoulli + untouched	83.576667	84.270000
4	Bernoulli + stretched	82.521667	83.690000
5	10 trees + 4 depth + untouched	69.641667	70.580000
6	10 trees + 4 depth + stretched	71.623333	73.170000
7	10 trees + 16 depth + untouched	99.071667	93.560000
8	10 trees + 16 depth + stretched	99.533333	94.430000
9	30 trees + 4 depth + untouched	74.141667	74.790000
10	30 trees + 4 depth + stretched	77.023333	78.310000
11	30 trees + 16 depth + untouched	99.405000	95.550000
12	30 trees + 16 depth + stretched	99.695000	96.410000

Mean images – Code on Page 11 - 12



Part 2 Code

Calculation of the Normal distribution parameters (Page 10 – 11)

```
1. for c in categories:
2.     XForC = X[Y == c]
3.     self.normDF[c] = {
4.         'mean': np.nanmean(XForC, axis=0),
5.         'var': np.nanvar(XForC, axis=0)
6.     }
7.     self.priors[c] = 1.0 * len(Y[Y == c]) / len(Y)
```

Calculation of the Bernoulli distribution parameters (Page 12 – 13)

```
1. self.priors[d] = log(1.0 * N / len(Y))
2. countOf1 = np.count_nonzero(XForD, axis=0)
3. countOf1 = (countOf1 + 1.) / (N + 2.) #Laplace smoothing
4. self.cond_probs[d] = countOf1
```

Calculation of the Naive Bayes predictions

For Gaussian (Page 11):

```
1. for c, g in self.normDF.items():
2.     mean, var = g['mean'], g['var']
3.     var = var + smoothing
4.     normPdf = norm.pdf(X, mean, np.sqrt(var))
5.     normPdf[normPdf == 0] = np.nan
6.     classConditionalProb = np.nansum(np.log(normPdf))
7.     P[c] = classConditionalProb + np.log(self.priors[c])
```

For Bernoulli (Page 13):

```
1. pred_class = None
2. max_ = float("-inf")
3. for d in self.priors:
4.     log_sum = self.priors[d]
5.     log_sum += np.sum(np.log(self.cond_probs[d][X == 1]))
6.     log_sum += np.sum(np.log(1 - self.cond_probs[d][X == 0]))
7.     if log_sum > max_:
8.         max_ = log_sum
9.     pred_class = d
```

Training of a decision tree (Page 14)

```
1. clf = RandomForestClassifier(n_estimators=numOfTrees, max_depth=maxDepth)
2. clf.fit(x_train, y_train)
```

Calculation of a decision tree predictions (Page 14)

```
1. y_pred_test = clf.predict(x_test)
2. correct_test = np.sum(y_pred_test == y_test)
3. accuracy_test = (correct_test / float(len(y_test))) * 100.0
```

Part A – complete code

```
import pandas as pd
import numpy as np
from scipy.stats import norm
import random
import math

raw_data = pd.read_csv('/Users/psaxena21/Documents/Mine/AML/pima-indians-diabetes.csv', header=None)
d = raw_data.values

class NaiveBayes:

    def __init__(self, ignoreMissingVal):
        self.ignoreMissingVal = ignoreMissingVal

    def testTrainSplit(self, data, ratio):
        localCopy = list(data)
        testSize = int(len(data) * ratio)
        testData = []
        while len(testData) < testSize:
            testData.append(localCopy.pop(random.randrange(len(localCopy))))
        testNPArr = np.array(testData)
        trainNPArr = np.array(localCopy)
        return trainNPArr[:, :8], trainNPArr[:, 8].astype(int), testNPArr[:, :8], testNPArr[:, 8].astype(int)

    def fit(self, X, Y):
        self.normDF = {}
        self.priors = {}
        categories = set(Y)
        if self.ignoreMissingVal:
```

```

X[X == 0] = np.nan

for c in categories:

    XForC = X[Y == c]

    self.normDF[c] = {

        'mean' : np.nanmean(XForC, axis=0),

        'var': np.nanvar(XForC, axis=0)

    }

    self.priors[c] = 1.0 * len(Y[Y == c]) / len(Y)

def __calculateProbability(self, x, mean, stdev):

    exponent = math.exp(-(math.pow(x-mean,2)/(2*math.pow(stdev,2))))

    return (1.0 / (math.sqrt(2*math.pi) * stdev)) * exponent

def __calculateLogNormPdf(self, X, mean, stddev):

    local_X = X[X != 0] if self.ignoreMissingVal else X

    local_mean = mean[X != 0] if self.ignoreMissingVal else mean

    local_stddev = stddev[X != 0] if self.ignoreMissingVal else stddev

    return np.log(np.array([self.__calculateProbability(local_X[i], local_mean[i], local_stddev[i]) for i in
range(len(local_X))]))

def predict(self, X):

    P = {}

    for c, g in self.normDF.items():

        # print "c:", c

        mean, var = g['mean'], g['var']

        classConditionalProb = 0

        # classConditionalProb = np.nansum(np.log(norm.pdf(X[X, mean, np.sqrt(var)))))

        classConditionalProb = np.sum(self.__calculateLogNormPdf(X, mean, np.sqrt(var)))

        P[c] = classConditionalProb + np.log(self.priors[c])

    bestCategory, bestProb = None, float("-inf")

    for category, probability in P.items():

```

```

        if bestCategory is None or probability > bestProb:

            bestProb = probability

            bestCategory = category

    return bestCategory

def runClassifier(data, ignoreMissingVal=False):

    accuracy = 0

    for i in range(10):

        nb = NaiveBayes(ignoreMissingVal)

        np.random.shuffle(data)

        X_Train, Y_Train, X_Test, Y_Test = nb.testTrainSplit(data, 0.20)

        nb.fit(X_Train, Y_Train)

        correct = 0

        for i in range(len(Y_Test)):

            Y_Pred = nb.predict(X_Test[i])

            if Y_Test[i] == Y_Pred:

                correct += 1

        accuracy += (correct/float(len(Y_Test)) * 100.0)

    return accuracy/10

acc = runClassifier(d)

print("Average accuracy over 10 test-train splits and without ignoring missing values is %f" % (acc))

acc = runClassifier(d, True)

print("Average accuracy over 10 test-train splits and ignoring missing values is %f" % (acc))

```

Part B – Complete code

```

%matplotlib inline

from matplotlib import pyplot as plt

import tensorflow as tf

from scipy.stats import norm

```



```

import numpy as np
import cv2
from math import log

mnist = tf.keras.datasets.mnist

(x_train, y_train), (x_test, y_test) = mnist.load_data()

ret, x_train_thresh = cv2.threshold(x_train, 127, 1, cv2.THRESH_BINARY)
ret, x_test_thresh = cv2.threshold(x_test, 127, 1, cv2.THRESH_BINARY)

def getStretchedImage(img):
    rows = np.any(img, axis=1)
    cols = np.any(img, axis=0)
    rmin, rmax = np.where(rows)[0][[0, -1]]
    cmin, cmax = np.where(cols)[0][[0, -1]]
    cropped = img[rmin:rmax, cmin:cmax]
    return cv2.resize(cropped, (20,20), interpolation=cv2.INTER_NEAREST)

def plotSampleFig(x_train_thresh, x_train_thresh_stretch):
    pixels = x_train_thresh[11232]
    img2 = cv2.resize(pixels, (20, 20), interpolation = cv2.INTER_NEAREST)
    img3 = getStretchedImage(pixels)
    f, axarr = plt.subplots(2,2)
    axarr[0,0].imshow(pixels, cmap=plt.cm.binary)
    axarr[0,1].imshow(img2, cmap=plt.cm.binary)
    axarr[1,0].imshow(img3, cmap=plt.cm.binary)
    axarr[1,1].imshow(x_train_thresh_stretch[11232], cmap=plt.cm.binary)
    plt.show()

x_train_thresh_stretch = np.array([getStretchedImage(x_train_thresh[i]) for i in range(len(x_train_thresh))])
x_test_thresh_stretch = np.array([getStretchedImage(x_test_thresh[i]) for i in range(len(x_test_thresh))])
plotSampleFig(x_train_thresh, x_train_thresh_stretch)

```

```

print("Shape of stretched training set " + repr(x_train_thresh_stretch.shape[0]))

x_train_flat_thresh = x_train_thresh.reshape(60000, 784)
x_test_flat_thresh = x_test_thresh.reshape(10000, 784)

x_train_flat_thresh_stretch = x_train_thresh_stretch.reshape(60000, 400)
x_test_flat_thresh_stretch = x_test_thresh_stretch.reshape(10000, 400)

def calculateAccuracy(nb, x_test, y_test, dataType="Untouched", setType="Test", modelType="Gaussian"):

    y_pred = np.apply_along_axis(nb.predict, 1, x_test)

    correct = np.sum(y_pred == y_test)

    accuracy = (correct/float(len(y_test)) * 100.0)

    print("Set-type: %s, DataType: %s, ModelType: %s, Accuracy: %f"%(setType, dataType, modelType, accuracy))

# Gaussian NB implementation
class NaiveBayesGaussian:

    def fit(self, X, Y):

        self.normDF = {}

        self.priors = {}

        categories = set(Y)

        for c in categories:

            XForC = X[Y == c]

            self.normDF[c] = {

                'mean': np.nanmean(XForC, axis=0),

                'var': np.nanvar(XForC, axis=0)

            }

            self.priors[c] = 1.0 * len(Y[Y == c]) / len(Y)

    def predict(self, X, smoothing=.01):

        P = {}

        for c, g in self.normDF.items():

            mean, var = g['mean'], g['var']

            var = var + smoothing

```

```

        normPdf = norm.pdf(X, mean, np.sqrt(var))

        normPdf[normPdf == 0] = np.nan

        classConditionalProb = np.nansum(np.log(normPdf))

        P[c] = classConditionalProb + np.log(self.priors[c])

        bestCategory, bestProb = None, float("-inf")

    for category, probability in P.items():

        if bestCategory is None or probability > bestProb:

            bestProb = probability

            bestCategory = category

    return bestCategory

nb = NaiveBayesGaussian()
nb.fit(x_train_flat_thresh, y_train)

print("Done training NB Gaussian untouched")

calculateAccuracy(nb, x_test_flat_thresh, y_test, "Untouched", "Test", "Gaussian")
calculateAccuracy(nb, x_train_flat_thresh, y_train, "Untouched", "Train", "Gaussian")


# Implementation for mean image plot

f, axarr = plt.subplots(5,2 , figsize=(15,15))

k = 0

mean_img_arr = [g['mean'] for c, g in nb.normDF.items()]

for i in range(5):

    for j in range(2):

        axarr[i, j].imshow(mean_img_arr[k].reshape((28,28)))

        k+= 1

plt.show()

nb = NaiveBayesGaussian()
nb.fit(x_train_flat_thresh_stretch, y_train)

print("Done training NB Gaussian stretched")

```

```
calculateAccuracy(nb, x_test_flat_thresh_stretch, y_test, "Stretched", "Test", "Gaussian")
calculateAccuracy(nb, x_train_flat_thresh_stretch, y_train, "Stretched", "Train", "Gaussian")
```

```
#Bernoulli NB implementation
```

```
class BernoulliNBClassifier(object):
```

```
    def __init__(self):
```

```
        self.priors = {}
```

```
        self.cond_probs = {}
```

```
    def fit(self, X, Y):
```

```
        digits = set(Y)
```

```
        for d in digits:
```

```
            XForD = X[Y == d]
```

```
            N = len(Y[Y == d])
```

```
            self.priors[d] = log(1.0 * N / len(Y))
```

```
        """Compute log( P(X|Y) )
```

```
        Use Laplace smoothing
```

```
        n1 + 1 / (n1 + n2 + 2)
```

```
        """
```

```
        countOf1 = np.count_nonzero(XForD, axis=0)
```

```
        countOf1 = (countOf1 + 1.) / (N + 2.)
```

```
        self.cond_probs[d] = countOf1
```

```
    def predict(self, X):
```

```
        """Make a prediction from text
```

```
        """
```

```
        pred_class = None
```

```
        max_ = float("-inf")
```

```

# Perform MAP estimation
for d in self.priors:
    log_sum = self.priors[d]
    log_sum += np.sum(np.log(self.cond_probs[d][X == 1]))
    log_sum += np.sum(np.log(1 - self.cond_probs[d][X == 0]))
    if log_sum > max_:
        max_ = log_sum
        pred_class = d

return pred_class

nb = BernoulliNBClassifier()
nb.fit(x_train_flat_thresh, y_train)
print("Done training NB Bernoulli untouched")
calculateAccuracy(nb, x_test_flat_thresh, y_test, "Untouched", "Test", "Bernoulli")
calculateAccuracy(nb, x_train_flat_thresh, y_train, "Untouched", "Train", "Bernoulli")

nb = BernoulliNBClassifier()
nb.fit(x_train_flat_thresh_stretch, y_train)
print("Done training NB Bernoulli stretched")
calculateAccuracy(nb, x_test_flat_thresh_stretch, y_test, "Stretched", "Test", "Bernoulli")
calculateAccuracy(nb, x_train_flat_thresh_stretch, y_train, "Stretched", "Train", "Bernoulli")

#DecisionTrees
from sklearn.ensemble import RandomForestClassifier

def runDecisionForestClassifier(x_train, y_train, x_test, y_test, numOfTrees, maxDepth, dataType="Untouched"):
    clf = RandomForestClassifier(n_estimators=numOfTrees, max_depth=maxDepth)
    clf.fit(x_train, y_train)
    y_pred_test = clf.predict(x_test)

```

```

y_pred_train = clf.predict(x_train)

correct_test = np.sum(y_pred_test == y_test)
correct_train = np.sum(y_pred_train == y_train)
accuracy_test = (correct_test/float(len(y_test)) * 100.0)
accuracy_train = (correct_train/float(len(y_train)) * 100.0)

print("%s image, %d trees, %d depth, test-set-accuracy: %f, train-set-accuracy: %f" %(dataType, numOfTrees,
maxDepth, accuracy_test, accuracy_train ))

runDecisionForestClassifier(x_train_flat_thresh, y_train, x_test_flat_thresh, y_test, 10, 4)
runDecisionForestClassifier(x_train_flat_thresh, y_train, x_test_flat_thresh, y_test, 30, 4)
runDecisionForestClassifier(x_train_flat_thresh, y_train, x_test_flat_thresh, y_test, 10, 16)
runDecisionForestClassifier(x_train_flat_thresh, y_train, x_test_flat_thresh, y_test, 30, 16)
strch = "Stretched"

runDecisionForestClassifier(x_train_flat_thresh_stretch, y_train, x_test_flat_thresh_stretch, y_test, 10, 4, strch)
runDecisionForestClassifier(x_train_flat_thresh_stretch, y_train, x_test_flat_thresh_stretch, y_test, 30, 4, strch)
runDecisionForestClassifier(x_train_flat_thresh_stretch, y_train, x_test_flat_thresh_stretch, y_test, 10, 16, strch)
runDecisionForestClassifier(x_train_flat_thresh_stretch, y_train, x_test_flat_thresh_stretch, y_test, 30, 16, strch)

```

References for code:

<https://machinelearningmastery.com/naive-bayes-classifier-scratch-python/>

<https://github.com/GPSingularity/Machine-Learning-in-Python/blob/master/nbsingularity.py>

https://mattshomepage.com/articles/2016/Jun/07/bernoulli_nb/

<https://nlp.stanford.edu/IR-book/html/htmledition/the-bernoulli-model-1.html>